

Attention Is All You Need · Transformer Architecture Benchmark

Abstract:

We propose the Transformer, a novel network architecture based solely on attention mechanisms, dispensing with

1. Introduction

Recurrent neural networks, particularly long short-term memory (LSTM) and gated recurrent (GRU) neural networks,

2. Methodology & Model Architecture

The Transformer follows an encoder-decoder structure using stacked self-attention and point-wise, fully connected

3. Empirical Results

On the WMT 2014 English-to-German translation task, the big transformer model achieves 28.4 BLEU, outperforming

4. Limitations & Future Work

Training the big model requires 8 V100 GPUs for 3.5 days. Future work includes extending the Transformer to process